

Hardware-oriented optimisation in CUDA

University of Regensburg



Tobias Sizmann

day month year

Abstract

Contents

1	Introduction	4
1.1	From single-core CPUs over multi-core CPUs to GPUs	4
1.2	What is CUDA	4
2	Tree Reduction on GPUs	7
2.1	The importance of reductions	7
2.2	The tree reduction algorithm	7
2.3	Naive implementation with CUDA	8
2.3.1	Serial tree reduction on a CPU	8
2.3.2	Tree reduction in CUDA for small arrays	10
2.3.3	Tree reduction in CUDA for arbitrary array sizes	12
2.4	Conclusion	14
3	Optimisations	15
3.1	Starting point: The naive kernel	15
3.2	Divergent warps	16
3.3	Memory bank conflicts	18
3.4	Idle threads after load	18
3.5	Automatic synchronisation within a warp	19
4	Benchmarks	21
A	Appendix One	22

Chapter 1

Introduction

1.1 From single-core CPUs over multi-core CPUs to GPUs

With the invention of the MOS transistor in 1959 and more specifically the silicon-gate MOS transistor in 1968 the development of single-chip microprocessors started. The first processors were single-threaded and their bottleneck was (besides other things) their clock rate. However, clock rates increased exponentially over time and subsequently did the computing power. In the 2000s this development was eventually brought to a halt at around 3.4 GHz due to thermodynamical limits of silicon. While pushing past that limit is possible, the extra costs for cooling usually outweigh the increase in computing power. This is the beginning of the multiprocessing era. The idea is simple: When one thread cannot run faster, just increase the number of threads. In 2006 the first desktop PCs with two cores were sold and ever since the number of threads is steadily increasing. There was one problem in particular, however, that CPUs (Central Processing Unit) were not efficient in solving, namely, rendering graphics. This task required a lot of independent and small calculations that needed to be done in real time - a prime example for a massively parallelisable computation. Since rendering graphics required such a different type of computing, the first GPUs (Graphics Processing Unit) were developed. These GPUs featured less single thread computation power but have a higher number of threads. State of the art GPUs have thousands of threads. This being said, the number of threads cannot be easily compared between a CPU and a GPU or even between different GPUs as they follow different design paradigms. How the GPU threads work exactly will be explored in this work. Generally, one can say that GPUs perform well in massively parallelisable computations where the single computations are not complex while CPUs excel at complicated single threaded problems (for example running the event loop of a desktop application). This shift to a higher number of threads rather than thread quality has been greatly motivated by scientific calculations, machine learning and graphics.

1.2 What is CUDA

Writing code for a GPU is not as straight forward as for a CPU since it is more dependent on its hardware. Different GPU developers use different application pro-




Figure 1.1: Comparison of the basic architectural differences of a GPU and a CPU. In green the random access memory (RAM), in yellow the cache, in blue the flow control unit and in red the arithmetic-logic unit (ALU). Both designs feature a RAM that all threads have access to. The GPU is split into many small "CPUs" (vertical groups) with their own flow control and cache. These are called streaming multiprocessors (SMPs or SMs). Generally the cache hierarchy is more complicated as depicted in the figure. However, the defining property here is that there exists a non-global cache level that is assigned to a group of ALUs, namely the SMs. Note that the notion of an ALU has slightly different meaning for a CPU and a GPU. For a CPU one ALU usually corresponds to one thread. For a GPU one ALU usually corresponds to a group of threads (for modern GPUs: 32), called a warp.

gramming interfaces (APIs). The biggest GPU designer Nvidia developed a framework called CUDA, which will be used in this work. It is simple to learn, offers efficient implementations and a plethora of literature and support. The downside is that it only supports Nvidia GPUs and is not open-source. An alternative would be OpenCL, which is open-source but harder to learn.

CUDA works as an extension to the C programming language and comes with its own compiler. Sections of code are distributed to either the CPU (called host) or the GPU (called device). When only writing code for the host normal C is used. Writing device code is more sophisticated. Here, the programmer first needs to write a so called kernel, which can be thought of the interior of a for-loop. Then the host must transfer required data to the memory of the GPU and call the kernel. When calling the kernel, the boundaries of the for-loop are set. The loop is then executed in parallel on the GPU. This API only allows parallelisation of for-loops. While this might seem like a very strict limitation, it closely relates to how a Nvidia GPU works.

The GPU can be thought of being organised in streaming multiprocessors (SMPs or SMs), warps and threads. Warps will be explained in more detail later. A SM groups together a set of (hardware) threads and allows synchronisation between them. Threads of different SMs cannot be synchronised. Also, threads within a SM share a cache, which cannot be accessed by threads of another SM. The domain of the for-loop is organised in blocks and (CUDA) threads, where the threads are grouped in blocks. The blocks are executed by the SMs, which means that only threads within a block can be synchronised and have access to the same cache. Note the differentiation of hardware and CUDA threads. This is an unfortunate ambiguity in the terminology of CUDA. While each CUDA thread is mapped to one hardware thread, a hardware thread can be mapped to (i.e. execute) several, none or one CUDA threads. This will become clearer in the next chapter.

Chapter 2

Tree Reduction on GPUs

2.1 The importance of reductions

A reduction in terms of parallel programming is an operation, where a large dataset of entries (e.g. numbers) is reduced to one entry. The simplest example is the sum of a set of numbers and will be used for the rest of this work. More generally, a reduction is defined by an operation $\circ : X \times X \rightarrow X$ on two entries with the following properties:

$$a \circ b = b \circ a \quad (\text{commutativity}), \quad (2.1)$$

$$a \circ (b \circ c) = (a \circ b) \circ c \quad (\text{associativity}). \quad (2.2)$$

This guarantees that the result is (mathematically) independent of the order in which the elements are reduced. Note, that these properties do not ensure numerical stability.

The reduction operation, first and foremost the sum, plays a crucial role in all of numerics. Simple linear algebra operations like the scalar product or the matrix multiplication already include a reduction: The entries of two vectors are multiplied element wise and the summed up. In machine learning, reductions are present as a key step in feed-forward networks: Again all the values of the node one layer below are multiplied by weights elementwise and the summed up to calculate the value of a singular node above. Reductions are a very basic and fundamental operation and, therefore, an efficient implementation is required.

2.2 The tree reduction algorithm

The most efficient algorithm is heavily dependent on the underlying hardware. For example, a node network with a lot of parallelisation overhead and a complicated communication topology might use a ring algorithm ("add value and pass to next"). In this work the focus is on single GPU reductions. Here, the tree reduction approach is the most successful. All available threads are called to reduce two entries to one in parallel, effectively halving the size of the dataset in one step. This is repeated until only one entry remains.




Figure 2.1: Example of a tree reduction of ten elements. Note how there is a leftover after the second reduction. These are usually handled by zeropadding (i.e. adding zeros) after each step to make the number of elements divisible by 2.

2.3 Naive implementation with CUDA

For a better understanding of the framework of CUDA and tree reduction algorithm an unoptimized code is presented which implements the algorithm for the case of addition of signed 32-bit integers. Note, that even though this is unoptimized GPU code, it runs magnitudes faster than on a CPU (exact speedup depending on the problem size and the available hardware).

When approaching a problem using CUDA the big challenge is to map the for-loop that one wants to parallelize to blocks and CUDA-threads. Often, the best starting point is to write serial CPU code to get a feeling for the problem and to have a working solution to test the optimized solutions against later. In this case, we will first implement a serial tree reduction in C and port this code to CUDA in a second step.

2.3.1 Serial tree reduction on a CPU

Our starting point is an integer array `h_in`. The `h_` denotes data stored on the host (CPU RAM), which is a useful convention, since CUDA does not distinguish between pointers to host and device memory. For simplicity, we assume that the size of the array is a power of two. If this is not the case, one could simply zeropad the data and would still be able to use the code.

Lets write a function `reduce` that executes the tree reduction. This function takes an input array `h_in`, performs the reduction and writes the result to an output `h_out`. We need two for-loops: One to iterate over the step of the reduction (vertically in fig. 2.3.1) and one to iterate over the (remaining) data set (horizontally). A temporary array is used to do the reduction on, so the input array stays




Figure 2.2: Sketch of the implementation of the tree reduction algorithm presented in this section. The cells denote entries in the data array. Two arrows pointing into a cell denote addition of the source cells values and writing of the result into the target cell. The final result is found in the first cell of the array.

untouched.

```

1 void reduce(int * h_in, int * h_out, int len) {
2     // copy data into temporary array
3     int * temp = (int *) malloc(len*sizeof(int));
4     memcpy(temp, h_in, len*sizeof(int));
5
6     int stride = 1;
7     // iterate over reduction steps (vertically)
8     for (int step = 0; step < log2(len); step++) {
9
10        // iterate over dataset (horizontally)
11        for (int i = 0; i < len; len += 2*stride) {
12            temp[i] = temp[i] + temp[i + stride];
13        }
14        stride *= 2;
15    }
16
17    // return the result
18    *h_out = temp[0];
19    free(temp);
20    return;
21 }

```

The inner loop of this implementation could be parallelised easily, since each step of the inner loop is independent. This means, however, that with each step of the outer loop all threads are created and destroyed, which leads to a large parallelisation overhead. In general and independent of CPU or GPU, one should always aim to parallelize the outermost loop. Here, the outer loop cannot be parallelized, since each step is dependent on the result of the step before. To fix this, one can swap the

two loops. While this requires slightly more code, it greatly simplifies all examples in the following. The rewritten function looks like this:

```

1 void reduce(int * h_in, int * h_out, int len) {
2     // copy data into temporary array
3     int * temp = (int *) malloc(len*sizeof(int));
4     memcpy(temp, h_in, len*sizeof(int));
5
6     // iterate over dataset (horizontally)
7     for (int i = 0; i < len; len += 1) {
8         int stride = 1;
9
10        // iterate over reduction steps (vertically)
11        for (int step = 0; step < log2(len); step++) {
12            if (i % (2 * stride) == 0) {
13                temp[i] = temp[i] + temp[i + stride];
14            }
15        }
16
17        stride *= 2;
18    }
19
20    // return the result
21    *h_out = temp[0];
22    free(temp);
23    return;
24 }

```

Note that an additional if-statement is required. This is a much better basis for parallelisation, since now the outer loop can be parallelised. However, one needs to be very careful as this is now prone to a race-condition. The parallelisation with CUDA is done in the next section.

It should be noted, that there are much faster ways to implement a reduction algorithm as a single-thread CPU application. This code has been written with the intent of parallelisation on a GPU and serves solely this purpose.

2.3.2 Tree reduction in CUDA for small arrays

Once the algorithm has been successfully implemented serially, the parallelisation of the target for-loop is always the same. One needs to map the domain of the for-loop (here integers from 0 to `len-1`) to threads. The maximum number of threads n_{Threads} a block can contain depends on the hardware, but is always a power of two. Modern Nvidia GPUs will allow for $n_{\text{Threads}} = 1024$, but it is not necessarily optimal to use all threads. This will be explored later. For now we will restrict the maximum size of our input array to 1024, such that only one block is required. With this, one can write down the kernel (terminology for a function running on the GPU).

```

1 __global__ void reduce(int * d_in, int * d_out, int len) {
2
3     // prepare shared data allocated by kernel invocation
4     // and copy input array
5     extern __shared__ int temp[];
6     temp[threadIdx.x] = d_in[threadIdx.x];
7     __syncthreads();
8
9     // do treereduction in interleaved addressing style

```

```

10     int stride = 1;
11     for (int step = 0; step < log2(len); step++) {
12
13         if (threadIdx.x % (2*stride) == 0) {
14             temp[threadIdx.x] += temp[threadIdx.x+stride];
15         }
16         __syncthreads();
17         stride *= 2;
18     }
19
20     // export result to global memory
21     if (threadIdx.x == 0) {
22         *d_out = temp[0];
23     }
24 }

```

There are several things to uncover here. First, the `__global__` void declares the function as a kernel, which runs on the device but can be invoked from the host. Secondly there is a new constant `threadIdx.x` available within the kernel. This is an identifier of the thread that is executing the kernel, is unique for all threads within a block and ranges from 0 to `numThreadsPerBlock - 1`. This replaces the index which the for-loop iterated over. The temporary array can be replaced by the ultra-fast cache shared by all threads of one block, which is allocated during kernel invocation and declared within the kernel by the line

```

1 extern __shared__ int temp[];

```

As mentioned earlier, there is a race condition between threads, which can be solved by the `__syncthreads()` method. This method acts as a wait-for-all barrier within a block (synchronisation between blocks is not possible!). Finally, the result needs to be exported. For this, only one thread is required, hence the if-statement. Note that the input and output pointer names start with a `d_`, which is not required but is a convention to mark, that these pointers live in the address-space of the device.

The kernel can be invoked from the host with the following code:

```

1 // allocate and copy to memory of device
2 int arraySize = 1024;
3 int * d_in, * d_out;
4 cudaMalloc(&d_in, sizeof(int)*arraySize);
5 cudaMalloc(&d_out, sizeof(int));
6 cudaMemcpy(d_in, h_in, sizeof(int)*arraySize,
7             cudaMemcpyHostToDevice);
8
9 // invoke kernel with the correct amount of threads and cache space
10 int numThreadsPerBlock = len;
11 int numBlocks = 1;
12 reduce <<< numBlocks, numThreadsPerBlock, sizeof(int)*
13             numThreadsPerBlock >>> (d_in, d_out, len);
14
15 // copy result to host and free memory
16 cudaMemcpy(h_out, d_out, sizeof(int), cudaMemcpyDeviceToHost);
17 cudaFree(d_in);
18 cudaFree(d_out);

```

First the data is copied to the device memory. Then the kernel is invoked by the triple angled brackets syntax. Within the brackets, the number of blocks n_{Blocks} (so far exactly one), the number of threads per block n_{Threads} and the amount of




Figure 2.3: Depicted is the schematic tree reduction in block structure for an array of length 16 and 4 threads per block. The first step requires 4 blocks and the second 1 block. The blocks run on the SMs of the GPU. The first step allows for four SMs and a total of 16 physical threads to be used in parallel. Between the two steps the output array is used as new input array. The second step can only make use of a single SM.

cache space are specified. The kernel arguments are specified in parenthesis. Finally the result is copied back to the host and the device memory is free'd. Note that in practice one would write more code to error-check every step, assure that the correct upstream is used and optimize the copy procedure. Nonetheless, this is a minimal working example.

The big problem with this implementation is, that it is limited to an array size of 1024, since we only use one block. This is solved in the next section.

2.3.3 Tree reduction in CUDA for arbitrary array sizes

The limitation to an array size n_{Data} of maximum block size can be solved in the following way: Split the array into equal chunks and run the kernel on each of these chunks individually. This procedure returns an array with size $n_{\text{Data}}/n_{\text{Threads}}$ (Assuming that n_{Data} is a power of 2, otherwise round up is required). The procedure is then again used recursively, until the array is reduced to a singular value. We refer to such a reduction step as "outer step", which reduces n_{Threads} elements, compared to an "inner step" reducing 2 elements.

While this could be implemented already with the existing kernel, there is a hardware structure available to do kernel invocations with a certain amount of threads in parallel, namely the streaming multiprocessors (SMs). The SMs are a second layer of parallelisation. One block of threads is always mapped to one SM, so several blocks can be executed in parallel on several SMs. Unlike the number of threads per block, the number of blocks is unlimited and blocks are queued until a SM is free to work on it. In our case this has the following implication. Lets assume that the initial array size is $n_{\text{Data}} = n \times 1024$. Then one could use n blocks of 1024 threads to do the first step of the outer reduction and use the full computational power of the GPU this way.

This leaves us with a new problem, however. Namely, how does one calculate the position of the array, one specific thread is supposed to work on? To this end,

CUDA offers two more constants in the kernel: `blockDim.x` and `blockIdx.x`. The first one is simply the number of threads in a block and the second one is a unique identifier of the current block, ranging from 0 to $n_{\text{Blocks}} - 1$. The array position can then be calculated with `threadIdx.x + blockIdx.x * blockDim.x`. The required amount of blocks for the kernel invocation can be calculated with rounding up division. This can be neatly done with $n_{\text{Blocks}} = (n_{\text{Data}} + n_{\text{Threads}} - 1) / n_{\text{Threads}}$, where the `"/` denotes integer division.

To enable the use of blocks with our kernel three small modifications need to be done. First, the access of our input array needs to be rewritten with the new formula for the position. Secondly, the output is not a singular value but an array. Thirdly, the range of the loop and the input array length can be deduced directly from the block size and number of threads per block and, therefore, the length of the input array does not need to be passed to the kernel. The result from i -th block should be written into the i -th position. Also for simplicity the loop over the `step`-variable has been replaced by a loop over `stride` directly, saving an extra variable and the `log2` function call. The finished code looks like this:

```

1 __global__ void reduce(int * d_in, int * d_out) {
2
3     // prepare shared data allocated by kernel invocation
4     // and copy input array
5     extern __shared__ int temp[];
6     temp[threadIdx.x] = d_in[threadIdx.x + blockIdx.x*blockDim.x];
7     __syncthreads();
8
9     // do treereduction in interleaved addressing style
10    for (int stride = 1; stride < blockDim.x; stride *= 2) {
11        if (threadIdx.x % (2*stride) == 0) {
12            temp[threadIdx.x] += temp[threadIdx.x+stride];
13        }
14        __syncthreads();
15    }
16
17    // export result to global memory
18    if (threadIdx.x == 0) {
19        *d_out = temp[blockIdx.x];
20    }
21 }
22 }

```

The invocation changes slightly. First, we need to specify the number of required blocks as argument in the triple angled brackets. Secondly, the kernel has only two arguments now. Also the size of the output array changes.

```

1 // set the number of threads per block and calculate the required
2 // number of blocks
3 int arraySize = sizeof(int)*len;
4 int numThreadsPerBlock = 1024;
5 int numBlocks = (arraySize + numThreadsPerBlock - 1) /
6                 numThreadsPerBlock;
7
8 // allocate and copy into device memory
9 int * d_in, * d_out;
10 cudaMalloc(&d_in, sizeof(int)*arraySize);
11 cudaMalloc(&d_out, sizeof(int)*numBlocks);
12 cudaMemcpy(d_in, h_in, sizeof(int)*arraySize,

```

```
11     cudaMemcpyHostToDevice);  
12 // invoke kernel with the correct amount of threads and cache space  
13 reduce <<< numBlocks, numThreadsPerBlock, sizeof(int)*  
14     numThreadsPerBlock >>> (d_in, d_out);  
15 // copy result to host and free memory  
16 cudaMemcpy(h_out, d_out, sizeof(int)*numBlocks,  
17     cudaMemcpyDeviceToHost);  
18 cudaFree(d_in);  
19 cudaFree(d_out);
```

Note that this code only executes a single outer step of the reduction. One would need to take the output array and feed it through this code until a single value is left. This was left out since it is mostly host code and not of particular interest for the rest of the work. It should be noted, however, that all code timings in the following were done for the full reduction.

2.4 Conclusion

This concludes the introduction to CUDA and tree reduction algorithms. From now on this work focuses on modifying the kernel to optimize it as much as possible. It should be noted, that a lot of optimization has been done already. For example, using the explicitly declared shared memory instead of DRAM or the implementation of the block structure already offer a great speedup compared to other algorithms one could come up with. Now, the main goal is to dive deeper into the hardware structure and achieve speedups by fixing problems like warp divergence and memory bank conflicts, which will be explained in full detail. Some algorithmic improvements will also be done and in the end all the optimisations will be benchmarked on several hardware systems.

Chapter 3

Optimisations

In this chapter various optimisations of the reduction kernel will be presented and discussed analogously to the talk by the author of [Har]. A reduction usually (at least in the case of addition) needs very few arithmetic operations and is, therefore, bottlenecked by memory access speed. This means that a good measure for performance is the achieved memory bandwidth rather than floating point operations per second. Especially, if one achieves bandwidths close to the maximum of the GPU, one can consider the kernel to be optimal. The bandwidth can be calculated using the formula

$$\text{bandwidth} = \frac{\text{input array size in bytes}}{\text{execution time}}. \quad (3.1)$$

The execution is defined as only the reduction and not the copying of the data to the GPU, which is reasonable, since for most use-cases the data would already be present in the GPU memory.

All achieved bandwidths shown in this chapter will be done on a Nvidia GeForce RTX 3070 (bandwidth: 448 GB/s) with an input array of $2^{27}(1.3 \cdot 10^8)$ 64-bit integers. This is a relatively new GPU (released Q4 2020). The GPU and the software that comes with it already have optimisations implemented that might mess with the results. To this end, data from an older setup (Nvidia G80) is provided as well published by [Har]. All measurements are done 1000 times to estimate the statistical error on the results. The results are presented merely to give a first impression on the effectiveness of the various changes on the kernel. In the last chapter, the performance will be investigated more closely and on several different GPUs.

3.1 Starting point: The naive kernel

The starting point is the kernel presented in the last chapter. In order to setup a timing routine, a wrapper was written, which executes the full reduction, i.e., it calls the kernel repeatedly with the required amount of blocks for each outer step until the reduction is fully done. Since the wrapper runs on the host, some of the execution time is spent on the CPU but for large enough arrays this part becomes negligible small. Another source of error is the call to the routine `cudaMemset()` which sets the memory of the device to a specified value. This is required for the aforementioned zeropadding. Again, the additional execution time is negligible. For the naive kernel a bandwidth of 127.1 ($\pm 1.0\%$) GBytes/s was measured.

3.2 Divergent warps

The first problem that needs to be dealt with is divergent warps. The threads of a streaming multiprocessor (SM) are clustered into so-called warps or SIMD (single instruction multiple data) lanes or vectors. Usually, one warp contains 32 threads and they are grouped with increasing `threadIdx.x`, i.e., threads 0 to 31 are one warp, 32 to 63 are one warp, etc. These threads are always synchronised in the sense that they execute the same instruction but act on different regions of memory (hence the name SIMD). If one thread in a warp were to execute a different instruction than the rest, which in CUDA happens through if-statements, all other threads are masked off (i.e. they still run but have no affect on memory). This is a highly inefficient procedure! Consider the following worst case example:

```

1 __global__ void horrific_kernel() {
2     if (threadIdx.x == 0) do_A();
3     if (threadIdx.x == 1) do_B();
4     if (threadIdx.x == 2) do_C();
5     if (threadIdx.x == 3) do_D();
6 }
```

Here, the first 4 threads all execute different paths, which is called divergent branching. In order to run this code the flow control unit of the SM must first mask off all threads of the warp except the first and run those instructions, then mask off all threads except the second and run those instructions and so on. This basically leads to the threads being executed serially rather than in parallel. There are two ways to fix this code. The first one is to use different blocks, i.e., use `blockIdx.x` in the if statements instead. This has two other disadvantages though. First, a whole SM is used to run the execution path of a single thread. Secondly, the threads cannot be synchronized afterwards. If one needs to add a barrier, basically a new and separate kernel invocation is required. The better solution is to use different warps:

```

1 __global__ void good_kernel() {
2     if (threadIdx.x == 0) do_A();
3     if (threadIdx.x == 32) do_B();
4     if (threadIdx.x == 64) do_C();
5     if (threadIdx.x == 96) do_D();
6 }
```

Each of the four paths are now running on different warps and therefore in parallel. This being said, there is still a lot of computer power wasted, since of each warp only one thread is used. Depending on the situation, however, this might be the most optimal solution. Usually, for highly complex branching code (e.g. an event loop of a desktop application), the CPU is preferred. Simple branching like in this example cannot be avoided most of the time and proper warp mapping is crucial for optimal performance.

In the case of the reduction kernel divergent warps are present which can be nicely seen in the code:

```

1 // do treereduction in interleaved addressing style
2 for (int stride = 1; stride < blockDim.x; stride *= 2) {
3     if (threadIdx.x % (2*stride) == 0) {
4         temp[threadIdx.x] += temp[threadIdx.x+stride];
5     }
6     __syncthreads();
7 }
```





Figure 3.1: Depicted is the same algorithm as in fig. 2.3.1 for 8 elements. However, this time the results of the addition of two cells is written into memory, such that the data stays contiguous. Note that one needs to be careful not to introduce an additional race condition when accessing the elements.

8 }

During the very first inner step, the if-statement branches within a warp. Thread 0 does some addition and thread 1 idles, thread 2 adds and thread 3 idles etc ... This means that there are two execution branches within one warp. If one would map all idle threads to one warp and working threads to another, then idling and adding would be executed in parallel. This can be simply done by rewriting the for-loop:

```

1 // do treereduction in interleaved addressing style
2 for (int stride = 1; stride < blockDim.x; stride *= 2) {
3     int index = 2 * stride * threadIdx.x;
4     if (index < blockDim.x) {
5         temp[index] += temp[index + stride];
6     }
7     __syncthreads();
8 }
```

Also, the costly %-operator vanishes this way.

The new bandwidth is 176.0 ($\pm 0.9\%$) GBytes/s, which is an increase of roughly 38%. In theory one would expect an increase close to 100%, and as a matter of fact, on older GPUs (Nvidia G80) with an older CUDA compiler, one actually achieves this. The reason here most likely is, that modern flow control units and compilers are able to do this optimisation to some degree by themselves, which means that the naive kernel already had some sort of divergent warp prevention. Still a 38% increase is non-negligible.

3.3 Memory bank conflicts

A similar problem to the divergent warps appears when accessing the ultra-fast shared cache. This memory is divided into so called memory banks (for modern Nvidia GPUs: 32), where each bank has a bandwidth of 32 bits per clock cycle. The mapping of the address space to the banks is done periodically, i.e.,

$$\text{bank number} = \text{address} \bmod 32 \quad (3.2)$$

Two threads cannot access the same bank in the same clock cycle. If this case appears, the warp(s) of the two threads is(are) frozen for one cycle and the memory loads/stores are done sequentially. One exception is, when the two threads load from the same address. In this case a broadcast operation is done, which requires no freezing.

Usually, memory bank conflicts are not completely avoidable, but must be reduced to a minimum. This can be achieved by memory coalescing, i.e., keeping the data that is being worked on contiguous in memory. In our case, the memory is not coalesced, which can be seen in fig. 2.3.1. After the first step there are "holes" in the data which is being accessed in the next step. This effectively halves the bandwidth of the shared memory in the second step. After the second step, it only gets worse. The bandwidth for the third step is effectively quartered. Since after each step the amount of data is also halved, this leads to an overall avoidable slow-down of roughly 50%.

The solution to this problem is simple. One needs to write the results of each addition back into memory in such a way, that the data stays coalesced. This is depicted in fig. 3.3. In terms of code this requires only a slight modification to the for-loop of the kernel:

```

1 // do treereduction in sequential addressing style
2 for (int stride = blockDim.x / 2; stride > 0; stride >>= 1) {
3     if (threadIdx.x < stride) {
4         temp[threadIdx.x] += temp[threadIdx.x + stride];
5     }
6     __syncthreads();
7 }
```

The order of the for-loop was reversed and the thread index can be used again for accessing the array.

With this a bandwidth of 185.1 ($\pm 1.1\%$) GBytes/s was achieved, which is an increase of roughly 5%. Again, on the older setup the expected speedup of 100% was achieved. The difference most likely stems from the compiler and flow control unit optimisations present in the newer setup.

3.4 Idle threads after load

While the last two optimisations were purely based on the architecture of the GPU, the next optimisation is of algorithmic character and more specific to tree reductions. The first step in the kernel is to load the data into the shared memory:

```

1 temp[threadIdx.x] = d_in[threadIdx.x + blockIdx.x*blockDim.x];
```

Here each thread of the kernel is active. However, after the load, i.e., in the first inner step, already half of the threads idle. One can make better use of these threads and optimise the first step, by combining loading and the first step of addition:

```
1 temp[threadIdx.x] =
2   d_in[threadIdx.x + 2*blockIdx.x*blockDim.x] +
3   d_in[threadIdx.x + 2*blockIdx.x*blockDim.x + blockDim.x];
```

One block of threads now loads double the amount of data, i.e., the data that was earlier assigned to two blocks. This means that the kernel needs to be invoked with only half the amount of blocks. The measured bandwidth is 346.4 ($\pm 0.4\%$) GBytes/s, which is an increase of roughly 87%. The older setup achieved a speedup of 78%. Since the compiler and the hardware are not able to make algorithmic improvements, it is expected that the speedups for the older and newer setup are in the same order of magnitude.

3.5 Automatic synchronisation within a warp

As mentioned before, threads within a warp are synchronized by hardware constraints if they do not branch. If they branch, threads of the same branch are still synchronized. This can be used to remove some `__syncthreads()` calls. In our case, if `stride \leq 32`, all non-idling threads are in a single warp. This means they are synchronized per se and no synchronisation barrier is required anymore. Additionally, for `stride < 32`, threads are divergent within one branch, which can be fixed by simply leaving the if-statement out in this case.

These changes are best implemented by writing a subroutine where the for-loop is unrolled explicitly for the case of `stride \leq 32`:

```
1 __device__ void warpReduce(volatile int * temp, int tIdx) {
2   temp[tIdx] += temp[tIdx + 32];
3   temp[tIdx] += temp[tIdx + 16];
4   temp[tIdx] += temp[tIdx + 8];
5   temp[tIdx] += temp[tIdx + 4];
6   temp[tIdx] += temp[tIdx + 2];
7   temp[tIdx] += temp[tIdx + 1];
8 }
```

The `__device__` statement declares the function to be a subroutine of the kernel. This function can only be called from the kernel. The `volatile` specifier is required to avoid compiler optimisations, that mess with the shared memory. Note that no if statements and no synchronisation barriers are required anymore.

The kernel itself changes to:

```
1 for (int stride = blockDim.x / 2; stride > 32; stride >>= 1) {
2   if (threadIdx.x < stride) {
3     temp[threadIdx.x] += temp[threadIdx.x + stride];
4   }
5   __syncthreads();
6 }
7
8 if (threadIdx.x < 32) warpReduce(temp, threadIdx.x);
```

The domain of the for-loop is shortened and the subroutine is called afterwards.

The achieved bandwidth is 413.0 ($\pm 0.7\%$) GBytes/s, which is an increase of roughly 20%. The older setup achieved a speedup of 80%. Unrolling for-loops is a classic compiler optimisation and most likely the reason for the discrepancy.

Chapter 4

Benchmarks

Appendix A

Appendix One

this will be space for the appendix

Bibliography

- [Har] Mark Harris. *Optimizing parallel reduction in CUDA*. URL: <https://developer.download.nvidia.com/assets/cuda/files/reduction.pdf>.